CSCI 570 - Fall 2022 - HW 5

September 26, 2022

Problem 1

Solve the following recurrences by giving tight Θ-notation bounds in terms of n for sufficiently large n. Assume that $T(\cdot)$ represents the running time of an algorithm, i.e. $T(n)$ is a positive and non-decreasing function of n. For each part below, briefly describe the steps along with the final answer.

(a) $T(n) = 4T(n/2) + n^2 \log n$

(b) $T(n) = 8T(n/6) + n \log n$

(c) $T(n) = \sqrt{6000}\, T(n/2) + n^{\sqrt{6000}}$

(d) $T(n) = 10T(n/2) + 2^n$

(e) $T(n) = 2T(\sqrt{n}) + \log_2 n$

Solution:

In some cases, we shall need to invoke the Master Theorem with one generalization as described next. If the recurrence $T(n) = aT(n/b) + f(n)$ is satisfied with $f(n) = \Theta(n^{\log_b a} \log^k n)$ for some $k \geq 0$, then $T(n) = \Theta(n^{\log_b a} \log^{k+1} n)$

(a) Observe that $f(n) = n^2 \log n$ and $n^{\log_2 4} = n^2$, so applying the generalized Master's theorem, $T(n) = \Theta(n^2 \log^2 n)$.

(b)Observe that $n^{\log_b a} = n^{\log_6 8}$ and $f(n) = n \log n = \mathcal{O}(n^{\log_6 8 - \epsilon})$ for any $0 < \epsilon < \log_6 8 - 1$. Thus, invoking the Master's theorem gives $T(n) = \Theta(n^{\log_b a}) = \Theta(n^{\log_6 8})$

(c) We have $n^{\log_b a} = n^{\log_2 \sqrt{6000}} = \mathcal{O}(n^{0.5 \log_2 8192}) = \mathcal{O}(n^{13/2})$. Further, $f(n) = n^{\sqrt{6000}} = \Omega(n^{70}) = \Omega(n^{13/2 + \epsilon})$ for any $0 < \epsilon < 63.5$. Also, $af(n/b) = \sqrt{6000}(n/2)^{\sqrt{6000}} \leq Cn^{\sqrt{6000}}$ holds for large enough n and some $C < 1$. Thus, from Master's theorem, $T(n) = \Theta(f(n)) = n^{\sqrt{6000}}$.

(d) We have $n^{\log_b a} = n^{\log_2 10}$ and $f(n) = 2^n = \Omega(n^{log_2 10+\epsilon})$ for any $\epsilon > 0$. Also, $af(n/b) = 10 \cdot 2^{n/2} \leq C \cdot 2^n$ holds for large enough n and some $C < 1$. Therefore, Master's Theorem implies that $T(n) = \Theta(f(n)) = \Theta(2^n)$.

(e) Use the change of variables: $n = 2^m$ to get $T(2^m) = 2T(2^{m/2}) + m$. Next, denoting $S(m) = T(2^m)$ implies that we have the recurrence $S(m) = 2S(m/2) + m$. Note that S(.) is a positive function due to the monotonicity of the increasing map $x \to 2^X$ and the positivity of $T(.)$. All conditions for applicability of Master's Theorem are satisfied and using the generalized version gives $S(m) = \Theta(m \log m)$ on observing that $f(m) = m$ and $m^{\log_2 2} = m$. We express the solution in terms of $T(n) = T(2^m) = S(m) = \Theta(m \log m) = \Theta(\log_2 n \log \log_2 n)$, for large enough n so that the expression is positive.

Rubric (15 pts)

- 3 pts: For each sub-part

  -1 pt: For the correct answer

  -2 pts: For the correct description

Problem 2

Solution:

In a set of cards, if more than half of the cards belong to a single user, we call the user a majority user.

Divide the set of cards into two roughly two equal halves, (that is one half is of size $\lfloor \frac{n}{2} \rfloor$ and the other, of size $\lceil \frac{n}{2} \rceil$). For each half, recursively solve the following problem, "decide if there exists a majority user and if she exists, find a card corresponding to her (as a representative)".

Once we have solved the problem for the two halves, we can combine them to solve the problem for the whole set, i.e. finding the global majority user. We can do that as follows.

If neither half has a majority user, then the whole set clearly does not have a majority user.

If both the halves have the same majority user, then that user is the global majority user. We can pick either one of the output cards returned by the halves as a representative for the whole set.

If the majority users are different, or if only one of them has a majority user, we need to check if any of these users is a global majority user. We can do this in a linear manner by comparing the representative card of the majority user with every other card in the whole set, counting the number of cards that belong to the same majority user.

If T(n) denotes the number of comparisons (invocations to the equivalence tester) of the resulting divide and conquer algorithm, then

$$T(n) \leq 2T(\lfloor \frac{n}{2} \rfloor) + \mathcal{O}(n) \rightarrow T(n) = \mathcal{O}(n \log n)$$

Rubric (15 pts)

• 7 pts: For describing the algorithm clearly

• 3 pts: For recurrence relation and time complexity computation

Problem 3

Solve Kleinberg and Tardos, Chapter 5, Exercise 5.

Solution:

Let $L = \{L_1, L_2, \ldots, L_n\}$ be the sequence of lines sorted in increasing order of slope. From now on, when we say sort a set of lines, it is in increasing order of slope. We divide the set of lines in half and solve recursively. When we are down to a set with only one line, we return the line as visible.

Recursively compute $L_{Bslash} = \{L_{i_1}, L_{i_2}, \ldots, L_{i_m}\}$, the sorted sequence of visible lines of the set $\{L_1, L_2, \ldots, L_{\lfloor \frac{n}{2} \rfloor}\}$. In addition compute the set of points $A = \{a_1, a_2, ..., a_{m-1}\}$ where a_j is the intersection of L_{i_k} and $L_{i_{k+1}}$.

Likewise compute $L_{slash} = \{L_{j_1}, L_{j_2}, \ldots, L_{j_r}\}$, the sorted sequence of visible lines of the set $\{L_{\lfloor \frac{n}{2} \rfloor + 1}, \ldots, L_n\}$. In addition compute the set of points $B = \{b_1, b_2, ..., b_{r-1}\}$ where b_j is the intersection of L_{j_k} and $L_{j_{k+1}}$.

Observe that by construction $\{a_1, a_2, ..., a_m\}$ and $\{b_1, b_2, ..., b_r\}$ are in increasing order of x-coordinate since if two visible lines intersect, the visible part of the line with smaller slope is to the left.

We now describe how the solutions for the two halves are combined. For this, we need to merge the two recursively computed sorted lists to get the list for the combined set of lines. The set of visible lines essentially forms a boundary, when seen from above. The intuition here is to find the point where the boundaries for the two halves intersect. This can then directly be used to find the boundary for the whole set, i.e. finding the

set of visible lines for the whole set. To locate the intersection point, we parse the two recursively-computed sorted lists to locate the first instance where a line from the first half is below a line from the second half. The intersection of these lines gives us the desired point. The figure below illustrates two examples of such intersections(Image Credits: https://cse.iitkgp.ac.in/ abhij/course/lab/Algo1/Spring20/A3.pdf)). This merging step can be done in $\mathcal{O}(n)$ time.

More specifically, we need to merge the two sorted lists A and B. Let $L_{up}(j)$ be the uppermost line in L_{Bslash}(j) and $\overline{L}_{up}$ the the uppermost line in L_{slash}. Let ℓ be the smallest index at which $\overline{L}_{up}$ is above L_{up}.

Let s and t be the indices such that $L_{up}(\ell) = L_{i_s}$ and define $\overline{L}_{up}(\ell) = L_{j_t}$

Let (a, b) be the intersection of $L_{up}(\ell)$ and $\overline{L}_{up}(\ell)$. This implies that $L_{up}(\ell)$ is visible immediately to the left of a and $\overline{L}_{up}(\ell)$ to the right. Hence the sorted set of visible lines of L is $L_{i_1}, L_{i_2}, \ldots, L_{i_{s-1}}, L_{i_s}, L_{j_t}, L_{j_{t+1}}, \ldots, L_r$.

The combination step takes $\mathcal{O}(n)$ time. If $T(n)$ denotes the running time of the algorithm, then

$$T(n) \leq 2T(\lfloor \frac{n}{2} \rfloor) + \mathcal{O}(n) \rightarrow T(n) = \mathcal{O}(n \log n)$$

Rubric (10 pts)

• 7 pts: For describing the algorithm clearly

• 3 pts: For recurrence relation and time complexity computation

Problem 4

Assume that you have a blackbox that can multiply two integers. Describe an algorithm that when given an n-bit positive integer a and an integer x, computes x^a with at most

$\mathcal{O}(n)$ calls to the blackbox.

Solution:

If a is odd:

$$x^a = x^{\lfloor \frac{a}{2} \rfloor} \times x^{\lfloor \frac{a}{2} \rfloor} \times x$$

If a is even:

$$x^a = x^{\lfloor \frac{a}{2} \rfloor} \times x^{\lfloor \frac{a}{2} \rfloor}$$

In either case, given $x^{\lfloor \frac{a}{2} \rfloor}$ it takes at most three calls to the black-box to compute x^a. We have thus reduced the problem of computing x^a to computing $x^{\lfloor \frac{a}{2} \rfloor}$ (which is an identical problem with input of size one bit smaller). Let $T(n)$ denote the running time of the corresponding divide-conquer algorithm. Thus

$$T(n) \leq T(n-1) + 3 \rightarrow T(n) = \mathcal{O}(n)$$

Remark: Such computations arise frequently in fields like cryptography. For instance, during the encryption of messages in the RSA cryptosystem, one has to make such a computation (m^e modulo a fixed number) where the exponent e is around 1024 bits. With the above algorithm, the number of multiplication calls to the blackbox is around 1024. If a naive method is used, it could take 2^{1024} multiplications which would be highly undesirable as this number far exceeds the number of atoms in the known universe!

Rubric (10 pts)

• 7 pts: For describing the algorithm clearly

• 3 pts: For recurrence relation and time complexity computation

Problem 5

Consider two strings a and b and we are interested in a special type of similarity called the "J-similarity". Two strings a and bare considered J-similar to each other in one of the following two cases: Case 1) a is equal to b, or Case 2) If we divide a into two substrings a_1 and a_2 of the same length, and divide bin the same way, then one of following holds: (a) a_1 is J-similar to b_1, and a_2 is J-similar to b^2 or (b) a_2 is J-similar to b_1, and a_1 is J-similar to b_2. Caution: the second case is not applied to strings of odd length.
Prove that only strings having the same length can be J-similar to each other. Further, design an algorithm to determine if two strings are J-similar within $O(n \log n)$ time (where n is the length of strings).

Solution:

To make our proof straightforward, we first change the statement to another equivalent: For every string a of length n and string b, b will be J-similar to a only if the length of b is equal to n. We use induction on n to prove the statement. The base case $(n = 1)$ is trivial. Assume we now have proved the statement for all $n < k$. Now we focus on $n = k$. Regarding to the definition, there are two cases that b can be J-similar to a. Firstly, if a is equal to b, then apparently they have the same length. On the other hand, if the second case happens, by the induction either one of the following will be correct: $\text{len}(a_1) = \text{len}(b_1)$, $\text{len}(a_2) = \text{len}(b_2)$ or $\text{len}(a_1) = \text{len}(b_2)$, $\text{len}(a_2) = \text{len}(b_1)$. In both case we have $\text{len}(a_1) + \text{len}(a_2) = \text{len}(b_1) + \text{len}(b_2)$, which means the length of b is also n.

The algorithm: we rearrange both of the two strings by some certain movements, then we prove they will project to the same string if and only if they are J-similar to each other. How to design this movement? To begin with, it's obvious that for any string a of even length, once we split it into two halves (a_1, a_2), $a' = a_2 a_1$ should be also J-similar to a. Based on this, we can design a divide and conquer approach to rearrange a, b into their lexicographically minimum equivalents. Specifically, we design a function J-sort(a) as follows:

```
function J-sort(a) {
    if len(a) % 2 == 1:
        return a # Unable to cut strings of odd length
    a₁, a₂ = a[:len(a)/2], a[len(a)/2:]  Cut string to half
    a₁ₘ = J-sort(a₁)
    a₂ₘ = J-sort(a₂)
    if a₁ₘ < a₂ₘ: # lexicographical order
        return a₁ₘ + a₂ₘ # concatenate
    else:
        return a₂ₘ + a₁ₘ
}
```

It's not hard to conclude that J-sort(a) has the lowest lexicographical order among all the strings that is J-similar to a (Just use induction like above). Because of this (and the same thing to b), we can further prove that a is J-similar to b if and only if J-sort(a) is equal to J-sort(b). Let $T(n)$ be the complexity of calculating J-sort(a) with respect to length n. Based on the previous analysis, we have $T(n) = 2 \cdot T(\frac{n}{2}) + O(n)$. The Master Theorem tells us $T(n) = O(n \log n)$. Aside from that, checking the equivalence between two strings costs only $O(n)$ time, which is negligible.

Rubric (15 pts)

- 5 points for proving the statement correctly

- 7 points for designing the Divide and Conquer Algorithm

Problem 6

Given an array of n distinct integers sorted in ascending order, we are interested in finding out if there is a Fixed Point in the array. Fixed Point in an array is an index i such that arr[i] is equal to i. Note that integers in the array can be negative

Example: Input: arr[] = -10, -5, 0, 3, 7 Output: 3, since arr[3] is 3

a) Present an algorithm that returns a Fixed Point if there are any present in the array, else returns -1. Your algorithm should run in $O(\log n)$ in the worst case.

b) Use the Master Method to verify that your solutions to part a) runs in $O(\log n)$ time.

c) Let's say you have found a Fixed Point P. Provide an algorithm that determines whether P is a unique Fixed Point. Your algorithm should run in $O(1)$ in the worst case

Solution:

a) First check whether the middle element is a Fixed Point or not. If it is, then return it; otherwise check whether the index of the middle element is greater than the value at the index. If the index is greater, then Fixed Point(s) lies on the right side of the middle point (obviously only if there is a Fixed Point). Else the Fixed Point(s) lies on the left side. If subproblem size equals one and Fixed is not found, return -1.

b) a=1, b=2, f(n) = O(1). $n^{log_b^a} = n^{log_2^1} = n^0 = O(1) \rightarrow$ Case 2: $T(n) = \theta(\log n)$

c) Say a fixed point is found at index i. Check indices i+1 and i-1. If we don't find fixed points at these two indices, then there cannot be any other fixed points since the array is sorted and all elements are distinct integers. Hence, for all other elements j above i, we must have arr[j] > j. And for all elements j below i, we must have arr[j] < j.

Rubric (12 pts)

• -3 if the incorrect half of array is pruned

• -1 if Master Theorem case is incorrect

• -1 if n**(log a base b) is incorrectly computed